

Activity 8 — From Hills to Backlog — Converting Ideas into Stories with GenAI

Course Fast-Track Innovations with Agile Design Thinking and Generative AI (GenAI)

WSQ Course Code TGS-2024049781

Topic Topic 03: Agile Development and AI for Rapid Solution Delivery

Duration 55 minutes · Group size: 3–5 participants

Tools Design Thinking Studio, Padlet Classroom Board, ChatGPT or Microsoft Copilot

Learning Outcome Addressed

LO3 · LO4 · K5 — Converting ideas into agile features, stories and tasks with GenAI; project management tools and techniques.

The Real Case

IBM runs a two-day Design Thinking and Agile workshop in which Day 1 produces 'Hills' — outcome statements in the form 'As a [who], I want [what], so that [wow]' — and Day 2 converts them into 20–30 user stories, prioritises them into three releases and sizes them with Planning Poker. In one documented session, a manufacturing user's need for inventory data started as a vague 'in minutes'. The team pushed until it became '5 minutes', and the justification made the requirement real: the data had to be available before the morning manager meeting. That specificity is what separates an outcome you can test from an aspiration you cannot. IBM prioritises releases as 'Cupcake, Birthday Cake, Wedding Cake' — where the cupcake is a complete, satisfying cake at small scale, not a slice of an unfinished one.

Your Scenario

Take the prototype concept your group validated in Activity 6. Your product trio — product manager, designer, engineer — now has one hour to turn it into a release-one backlog that a delivery team could actually start on Monday.

What You Will Do

Learners write a measurable Hill, use GenAI to expand it into user stories with acceptance criteria, critique and correct the AI output, prioritise into the three-cake release model and size the stories with relative estimation.

What you will produce

A prioritised release-one backlog: one Hill, 8–12 user stories with acceptance criteria, a Cupcake/Birthday/Wedding split and story-point estimates.

Step-by-Step Instructions

1. Open the Padlet classroom board, Activity 8 section.
2. Write ONE Hill for your concept: 'As a [who], I want [what], so that [wow]'. Make the wow measurable.
3. Post the Hill and have another group challenge its measurability. Revise it.
4. Prompt the AI: 'Convert this Hill into 10 user stories in the form As a/I want/so that, each with 2-3 testable acceptance criteria.'
5. Review every story. Delete the ones that assume something you never validated. Rewrite the weak acceptance criteria.
6. Split the surviving stories into Cupcake, Birthday Cake and Wedding Cake releases.
7. Size the Cupcake stories using relative points (1, 2, 3, 5, 8) — discuss any estimate where the group differs by more than two cards.
8. Post your final release-one backlog to Padlet and Like the backlog you would most want to inherit as a developer.

Discussion Questions

1. Compare a Hill ('so that [wow]') with a user story ('so that [reason]'). What does the Hill capture that the story loses?
2. The team pushed 'in minutes' to '5 minutes'. What forced the specificity, and what would have happened without it?
3. Review the AI-generated stories. What did it get structurally right, and what did it get wrong about your context?
4. Explain 'Cupcake' to someone who thinks an MVP is 'version one with features removed'. Where is the difference?
5. Story points measure size and complexity, not duration. Why does the distinction matter for a team's velocity?

How You Know You Are Done

Your Cupcake release is genuinely shippable and useful on its own, every story has testable acceptance criteria, and you can point to which stories you deleted from the AI output and say why.

Debrief — What the Room Should Conclude

GenAI is genuinely good at the mechanical layer of this task — it produces well-formed stories with plausible acceptance criteria in seconds, and it rarely forgets the syntax. What it cannot supply is context: your constraints, your regulatory environment, your legacy integration, your users' actual tolerance for change. The reliable pattern

is AI drafts, humans decide. Watch for the two characteristic AI failures in the room: stories that are technically well-formed but describe features nobody validated, and acceptance criteria that restate the story instead of defining a testable condition. On the cake model — the point of the cupcake is that a person could genuinely enjoy it and it stands alone. A slice of unfinished wedding cake is not a smaller product, it is a broken one, and shipping it teaches you nothing except that users dislike broken things.